

Exercises on Session 2: Probabilities and Random Variables

Vasilis Gkolemis

June 2025

Instructions These exercises are designed to help you practice probabilities and statistics.

1 Exercise: The expected value of the mean (independent case)

Let X_1, X_2, \dots, X_n be independent random variables. Prove that $Y = \frac{1}{n} \sum_{i=1}^n X_i$ has expected value:

$$\mathbb{E}[Y] = \frac{1}{n} \sum_{i=1}^n \mathbb{E}[X_i].$$

What is the expected value of Y if all X_i are coming from the same distribution with expected value μ ?

Solution.

$$\begin{aligned} \mathbb{E}[Y] &= \mathbb{E} \left[\frac{1}{n} \sum_{i=1}^n X_i \right] \\ &= \frac{1}{n} \mathbb{E} \left[\sum_{i=1}^n X_i \right] \\ &= \frac{1}{n} \sum_{i=1}^n \mathbb{E}[X_i]. \end{aligned}$$

If all X_i are coming from the same distribution with expected value μ , then:

$$\mathbb{E}[Y] = \frac{1}{n} \sum_{i=1}^n \mu = \mu.$$

2 Exercise: The variance of the mean (independent case)

Let X_1, X_2, \dots, X_n be independent random variables with variances $\sigma_i^2 = \text{Var}(X_i)$. Prove that the variance of the mean $Y = \frac{1}{n} \sum_{i=1}^n X_i$ is given by:

$$\text{Var}(Y) = \frac{1}{n^2} \sum_{i=1}^n \sigma_i^2.$$

What is the variance of Y if all X_i are coming from the same distribution with variance σ^2 ?

Solution. By the properties of variance and independence, we have:

$$\begin{aligned}\text{Var}(Y) &= \text{Var}\left(\frac{1}{n} \sum_{i=1}^n X_i\right) \\ &= \frac{1}{n^2} \text{Var}\left(\sum_{i=1}^n X_i\right) \\ &= \frac{1}{n^2} \sum_{i=1}^n \text{Var}(X_i) \\ &= \frac{1}{n^2} \sum_{i=1}^n \sigma_i^2.\end{aligned}$$

If all X_i are coming from the same distribution with variance σ^2 , then:

$$\text{Var}(Y) = \frac{1}{n^2} \sum_{i=1}^n \sigma^2 = \frac{n\sigma^2}{n^2} = \frac{\sigma^2}{n}.$$

3 Exercise: Practical implication of (1) and (2)

You are given a test and you are asked for the expected value of some probability distributions; the gamma distribution $\Gamma(\alpha, \beta)$ and the beta distribution $\text{Beta}(\alpha, \beta)$. Unfortunately, you do not remember the formulas for the expected values of these distributions, but you have access to a computer with R installed. How could you estimate the expected values of these distributions using sampling?

1. Write an R script to estimate the expected value of a Gamma distribution with parameters α and β (shape and rate).
2. Write an R script to estimate the expected value of a Beta distribution with parameters α and β .

Hint: In R, you can use the ‘`rgamma`’ and ‘`rbeta`’ to sample from the Gamma and Beta distributions, respectively. For ‘`rgamma`’, the function signature is ‘`rgamma(n, shape, rate)`’ where ‘`n`’ is the number of samples, ‘`shape`’ is α , and ‘`rate`’ is β . For ‘`rbeta`’, the function signature is ‘`rbeta(n, shape1, shape2)`’ where ‘`n`’ is the number of samples, ‘`shape1`’ is α , and ‘`shape2`’ is β .

After you have run the R script . Your solution will not be exact but it will be close enough to the expected value. Can you answer why it is close enough?

```
# 1. Estimate expected value of a Gamma(alpha, beta) distribution
alpha <- 2
beta <- 3 # Rate parameter
set.seed(123)
samples_gamma <- rgamma(10000, shape = alpha, rate = beta)
expected_gamma <- mean(samples_gamma)
cat("Estimated expected value of Gamma(", alpha, ",", beta, ") = ", expected_gamma, "\n")

# 2. Estimate expected value of a Beta(alpha, beta) distribution
alpha_beta <- 2
beta_beta <- 5
```

```

set.seed(123)
samples_beta <- rbeta(10000, shape1 = alpha_beta, shape2 = beta_beta)
expected_beta <- mean(samples_beta)
cat("Estimated expected value of Beta(", alpha_beta, ",", beta_beta, ") = ", expected_beta,

```

It will be close enough because from (1):

$$\mathbb{E}[Y] = \mu$$

and from (2):

$$\text{Var}(Y) = \frac{\sigma^2}{n}$$

So, for large n , the sample mean Y will converge to the expected value μ due to the Law of Large Numbers. The variance decreases as n increases, making the estimate more precise.

4 Exercise: Sample from a more complicated distribution

You have to sample from $p(x, y) = p(y|x)q(x)$, where $p(y|x)$ is a Gaussian distribution with mean $\mu = x$ and variance σ^2 , and $q(x)$ is a uniform distribution on the interval $[a, b]$. Can you do that in R? Is it straightforward (like two or three lines of code) or do you need to write a complicated script?

```

# Sample from p(x,y) = p(y|x)q(x)
set.seed(123)
a <- 0
b <- 10
n_samples <- 10000
# Sample x from uniform distribution
x_samples <- runif(n_samples, min = a, max = b)
# Sample y from Gaussian distribution with mean = x and variance = sigma^2
sigma <- 1
y_samples <- rnorm(n_samples, mean = x_samples, sd = sigma)
# Combine samples into a data frame
samples <- data.frame(x = x_samples, y = y_samples)
# Display first few samples
head(samples)

```

It is straightforward to sample from the joint distribution $p(x, y)$ in R. You first sample x from the uniform distribution and then sample y from the Gaussian distribution with mean equal to the sampled x and a fixed variance σ^2 .